

CISC 372

Advanced Data Analytics

L8 –Trees

	name	age	state	num_children	num_pets
0	john	23	iowa	2	0
1	mary	78	dc	2	4
2	peter	22	california	0	0
3	jeff	19	texas	1	5
4	bill	45	washington	2	0
5	lisa	33	dc	1	0



wild DATAFRAME appeared!

Monday

- Accuracy/Efficiency/Scalability/Interpretability
- Decision Tree
 - Top-down divide-and-conquer algorithm
 - Three equations:
 - Entropy
$$Info(D) = - \sum_{i=1}^m p_i \log_2(p_i)$$
 - Conditional Entropy
$$Info_A(D) = \sum_{j=1}^v \frac{|D_j|}{|D|} \times Info(D_j)$$
 - Information Gain
$$Gain(A) = Info(D) - Info_A(D)$$
 - Interpretability
 - Overfitting & Pruning


Today

- Gain Ratio
- Gini Index
- ID3, CART, C4.5
- Feature Selection (is difficult)
- Random Forest

Information Gain

- Assume that the attributes below that have the **SAME information gain**
 - A: age
 - B: has_1000_driving_tickets
 - C: student_id
 - Which one would you pick?

Split for ID Code Attribute



Entropy of split = 0 (since each leaf node is “pure”, having only one case).

Information gain is maximal for ID code

Gain ratio

- *Gain ratio*: a modification of the information gain that reduces its bias towards high-branch attributes
- Gain ratio should encourage
 - Even distribution (of instances/samples)
 - Low distinct values
- How?
 - Low intrinsic information => the entropy of the attribute itself

Gain Ratio and Intrinsic Info.

- Intrinsic information: entropy of distribution of instances into branches

$$\text{Intrinsic}(A) = - \sum_A^i \frac{|S_i|}{|S|} \log_2 \frac{|S_i|}{|S|}$$

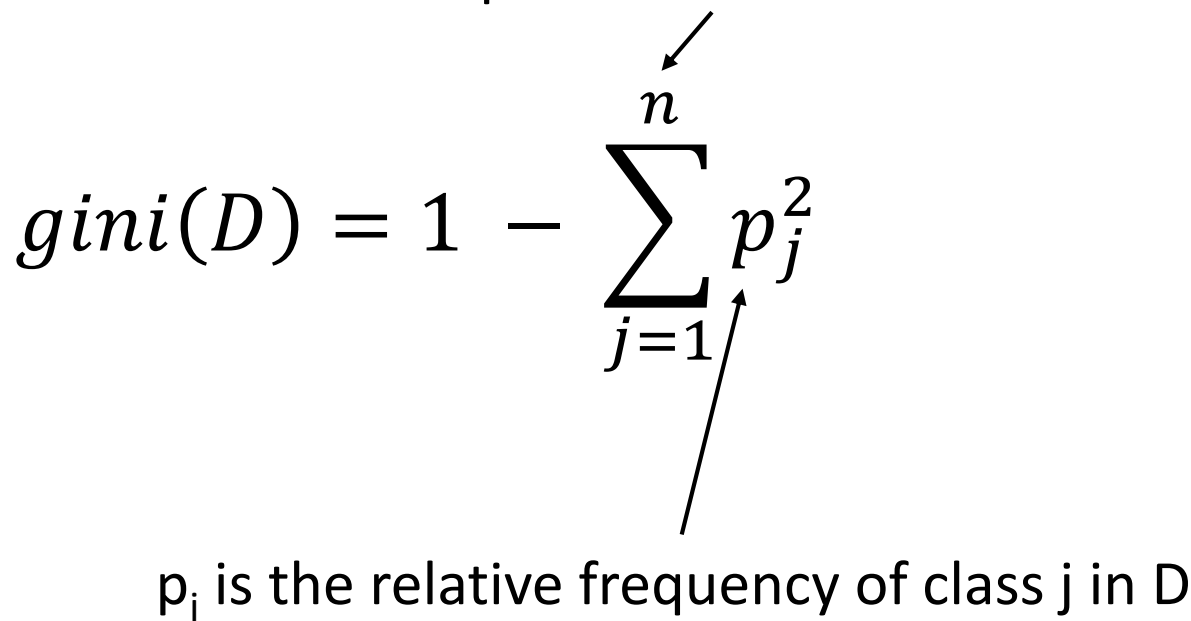
$$\begin{aligned} \text{GainRatio}(A) &= \frac{\text{Gain}(A)}{\text{Intrinsic}(A)} \\ &= \frac{\text{Info}(D) - \text{Info}_A(D)}{\text{Intrinsic}(A)} \end{aligned}$$

More on the gain ratio

- Problem with gain ratio: it may **overcompensate**
 - May choose an attribute just because its intrinsic information is very low
- Standard fix:
 - First, only consider attributes with greater than average information gain
 - Then, compare them on gain ratio

Gini Index - CART

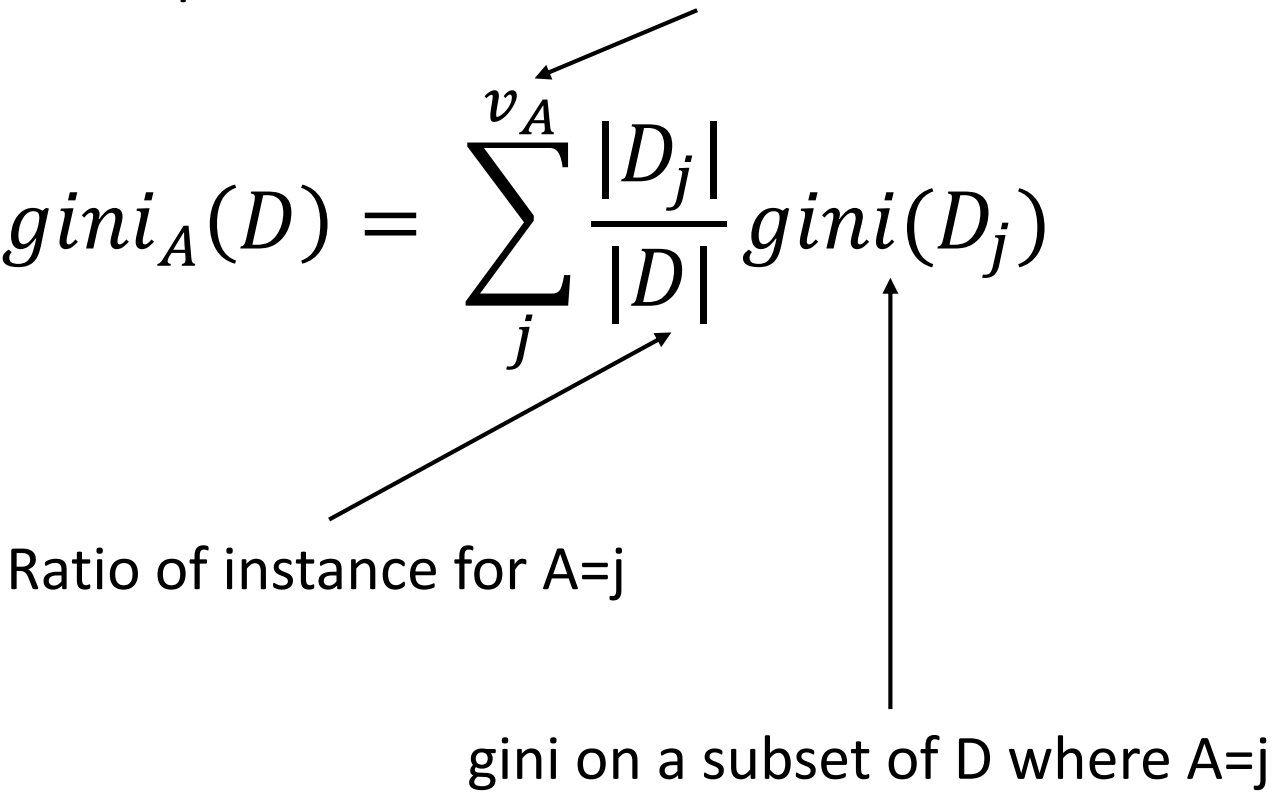
Number of unique class labels

$$gini(D) = 1 - \sum_{j=1}^n p_j^2$$


p_j is the relative frequency of class j in D

Gini Index - CART

Unique values for attribute A

$$gini_A(D) = \sum_j^{v_A} \frac{|D_j|}{|D|} gini(D_j)$$


Ratio of instance for A=j

gini on a subset of D where A=j

Gini Index - CART

- Algorithm for top-down induction of decision trees ("ID3") was developed by Ross Quinlan
 - C4.5 (gain ratio), which can deal with numeric attributes, missing values, and noisy data
- Similar approach: CART
- There are many other attribute selection criteria!

Numeric attributes

- Standard method: discretization
 - E.g. $\text{temp} < 45$
- Unlike nominal attributes, every attribute has many possible split points
- Solution is straightforward extension:
 - Evaluate info gain (or other measure) for every possible split point of attribute
 - Choose "best" split point
 - Info gain for best split point is info gain for attribute
- Computationally more demanding

Numeric attributes

- Standard method: discretization
 - E.g. $\text{temp} < 45$
- Unlike nominal attributes, every attribute has many possible split points
- Solution is straightforward extension:
 - Evaluate info gain (or other measure) for every possible split point of attribute
 - Choose "best" split point
 - Info gain for best split point is info gain for attribute
- Computationally more demanding

Prepruning

- Based on statistical significance test
 - Stop growing the tree when there is *no statistically significant* association between any attribute and the class at a particular node
- Most popular test: *chi-squared test*
- ID3 used chi-squared test in addition to information gain
 - Only statistically significant attributes were allowed to be selected by information gain procedure

Early stopping

	a	b	class
1	0	0	0
2	0	1	1
3	1	0	1
4	1	1	0

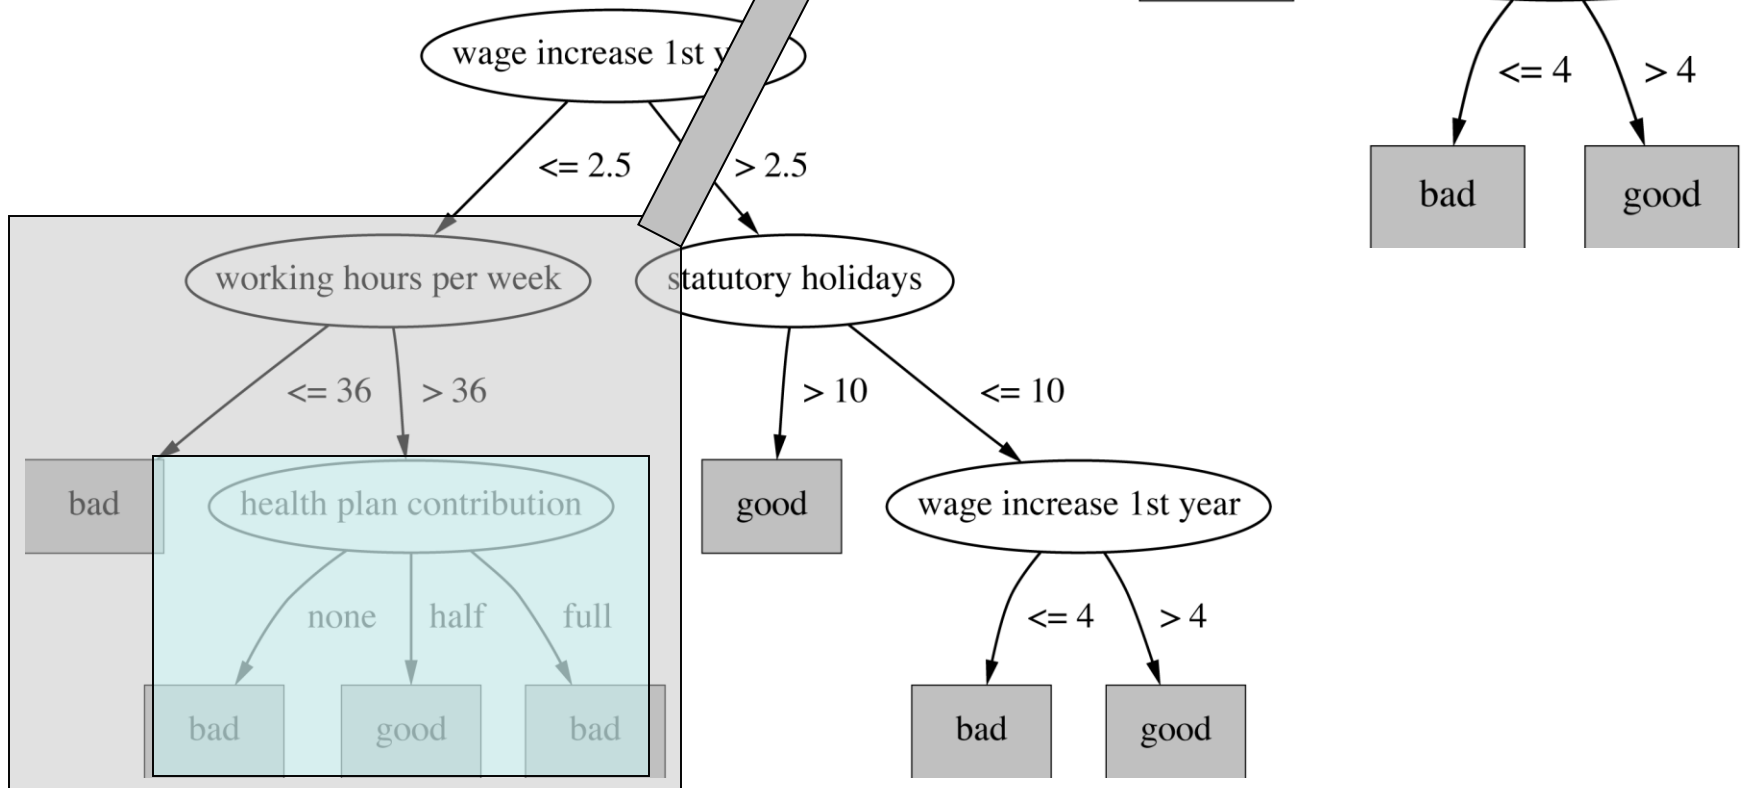
- Pre-pruning may stop the growth process prematurely:
early stopping
- Classic example: XOR/Parity-problem
 - No *individual* attribute exhibits any significant association to the class
 - Structure is only visible in fully expanded tree
 - Pre-pruning won't expand the root node
- But: XOR-type problems rare in practice
- And: pre-pruning faster than post-pruning

Post-pruning


- First, build full tree
- Then, prune it
 - Fully-grown tree shows all attribute interactions
- Problem: some subtrees might be due to chance effects
- Two pruning operations:
 1. *Subtree replacement*
 2. *Subtree raising*
- Possible strategies:
 - error estimation
 - significance testing
 - Minimum Description Length principle

Subtree replacement


- *Bottom-up*
- Consider replacing a tree only after considering all its subtrees




Subtree raising




- Delete node
 - Redistribute instances
 - Slower than subtree replacement
- (Worthwhile?)*



Decision Boundary



Decision Boundary – overfit (no pruning)



Feature Selection (is difficult)

- Why?
 - Given m features: high training complexity often m^2
- What (to remove)?
 - Values are not correlated with the target label
 - Correlation – no transitive property.
 - A, B can both correlate to D but A & B may/may not be correlated (so we can't just remove A/B)
 - A correlated to B \Rightarrow we can't just remove A
 - A, B correlates to D, and A, B correlates to each other \Rightarrow redundant
 - But keep A or B?
 - Theoretically we need to test every possible pair of m .
 - Low predictive power
 - A set of feature combined \geq sum of each individual
 - Need to evaluate every possible subset of m .

Feature Selection (is difficult)

- Why?
 - Given m features: high training complexity often m^2
- What (to remove)?
 - Values are not correlated with the target label
 - Correlation – no transitive property.
 - A, B can both correlate to D but A & B may/may not be correlated (so we can't just remove A/B)
 - A correlated to B \Rightarrow we can't just remove A
 - A, B correlates to D, and A, B correlates to each other \Rightarrow redundant
 - But keep A or B?
 - Theoretically we need to test every possible pair of m .
 - Low predictive power
 - A set of feature combined \geq sum of each individual
 - Need to evaluate every possible subset of m .

Feature Selection



Random Forests

The name comes from the fact that this was originally an extension of decision trees, but the idea applies to **any** predictor.

The basic idea: build a large number of decision trees, each using only some of the attributes of the dataset. Deploy the forest using voting (classification) or averaging (regression).

Recall: Bootstrapping

Recall: a bootstrap sample from n objects means draw n objects with replacement.

The sample will contain about $2n/3$ unique objects, leaving $n/3$ objects that were never selected.

These unselected objects can be used as a test set with nice properties.



Growing one tree of the forest:

1. Select a bootstrap sample of the objects^s
2. If there are m attributes, choose some $k < m$ (usually k is much smaller than m , maybe \sqrt{m})
3. Grow a decision tree by randomly selecting k of the m attributes at each level, and choosing the best split based on these k attributes (in whatever standard way)

4. Grow the decision tree to full size (i.e. no pruning)
5. Use the **out-of-bag set** as a **test sample** and measure the test error for this tree (ongoing estimate of how well we're doing)
6. Run the entire dataset through the tree. Whenever **two objects end up at the same leaf** increase the **proximity** score for the pair.

Growing the forest:

1. Grow some number of individual trees – this is quite cheap so we might grow 100s or 1000s.
2. For each object, consider how many times it appeared in a test set, and count

$$\frac{\text{number of times prediction was wrong}}{\text{number of times it appeared in the test set}}$$

Average this to get an overall test error

3. Divide the proximities by the total number of trees

Advantages:

- handles multiple classes
- fast (worse than 1 decision tree, better than a neural network)
- low variance, low bias
- no separate test procedure needed (no cross-validation)
- high accuracy
- effective for large numbers of attributes
- helpful for attribute selection
- does not overfit no matter how many trees

Disadvantages:

- opaque predictor
- expensive to deploy